

HelixCode — Fixed Items Summary

Generated from docs/Fixed.md per Constitution §11.4.19.
Counts only.

Round 189 prefix convention: Open- and closed-item IDs are now scope-prefixed (HXC, HXA, HXL, HXQ, VEN, ...). See docs/Issues.md “Prefix convention” section for the full table + legacy ISSUE-NNN mapping. The aggregate counts below are unaffected by the rename — only labels change.

Aggregate counts (post round-system rebaseline 2026-05-19)

Type	Count	Closure vocabulary (CONST-057)
Bug	19	Fixed (→ Fixed.md)
Feature	68	Implemented (→ Fixed.md)
Task	7	Completed (→ Fixed.md)

Total closed items: 96 (in the round-system tracker; pre-round closures tracked separately in docs/improvements/PROGRESS.md). Round 463 added HXC-003 (CONST-046 i18n migration backlog — Implemented (→ Fixed.md): the CONST-046 i18n migration campaign is concluded. The genuine user-facing (C) string-literal surface — UI text, prompts shown to the operator, error messages surfaced to end users, labels, helper text — is exhausted across all 7 CONST-046 scope areas: helix_code internal/ + cmd/ + applications/ (exhausted rounds 461/462), LLMsVerifier (round 452), helix_qa, and all owned vasic-digital/* + HelixDevelopment/* submodules (rounds 413/441). Across ~91-462 rounds tens of thousands of literals were migrated through nicksnyder/go-i18n/v2 Bundle/Localizer seams plus locale-aware resource files, and every migration round shipped paired-mutation anti-bluff tests that plant a known un-migrated literal and assert the audit-gate reports FAIL — so a PASS certifies real i18n coverage. The audit-gate scripts/audit-const046-hardcoded-content.sh --fail-on-new is enforced and each round re-ran --update-baseline so the snapshot shrank monotonically. The remaining ~55k audit-baseline hits are all OUT of CONST-046 scope — LLM prompt templates, wrapped-error developer-facing tech strings, identifier tokens, struct-tag

keys, format-spec tokens, test fixtures — classified file-by-file in docs/audits/2026-05-20-internal-const046-classification.md. CONST-046's invariant is satisfied: no user-facing text is hardcoded as a static literal). Round 403 added HXC-008 (CONST-055 G1 governance gaps surfaced by the post-constitution-pull validation sweep — two pre-existing cascade-drift defects: (a) docs/improvements/submodule_owned.txt line 10 referenced a non-existent dependencies/HelixDevelopment/Models path which the CONST-052 batch-1 rename alea3c8 had lowercased to models, corrected to dependencies/HelixDevelopment/models — verify-governance-cascade.sh re-run dropped the false path does not exist on disk FAIL; (b) helix_qa/CONSTITUTION.md was missing anchors CONST-047..057 and VisionEngine/CONSTITUTION.md was missing §11.4.69 — both CLAUDE.md/AGENTS.md already carried them, so the CONST-047..057 block was cascaded into helix_qa/CONSTITUTION.md from the meta CONSTITUTION.md and §11.4.69 into VisionEngine/CONSTITUTION.md; verify-governance-cascade.sh → 0 failures, verify-all-constitution-rules.sh → 6 gates / 0 failures; helix_qa 1364d23 + VisionEngine b3a13d8 pushed to all upstreams), HXC-007 (Constitution §11.4.68/70-74 cascade verified genuinely complete — constitution 584b3ee→34a82b3, 6 rules cascaded to the meta-repo + 67 owned submodules, meta .gitmodules pointer confirmed at 34a82b3), and HXC-009 (owned-submodule GitHub↔GitLab mirror-divergence reconciliation verified complete — helix_qa, VisionEngine, LLMProvider, challenges, containers, DocProcessor all reconciled via merge-first per CONST-061/§11.4.71, all converged + pushed to all upstreams, meta pointers bumped 68309b8e). Round 402 added HXC-011 (helix_qa runner emitted hollow sub-microsecond PASSED rows for desktop-platform bank cases instead of executing the case's action: — §11.4 / CONST-035 PASS-bluff in the QA runner itself; reproduce-before-fix Rule 7 / §11.4.43 TDD-fix; root cause: definitionChallenge.Execute unconditionally returned Status=Skipped and never shelled out, and the generic challenges/pkg/bank loader drops each case's steps array so executable action data never reached the orchestrator; RED-with-fix-reverted captured the bluff — a desktop shell: sentinel-write action ran through helixqa run -platform desktop and the sentinel was never written + a shell: exit 17 case produced 0 failures; fix all in helix_qa keeping the challenges submodule decoupled per CONST-051(B) — new testbank.ActionTypeShell runs a host command via os/exec, loadExecutableCases() re-parses banks through pkg/testbank which preserves steps, runPlatform builds a per-platform registry of definitionChallenge wrappers carrying the executable case + target platform, and Execute on desktop runs each shell: step capturing the real exit code + output and scores PASS only when every step exits 0 — non-zero scores FAIL, no-executable-action cases emit an honest SKIP with explicit reason per §11.4.3; GREEN: 3 HXC-011 tests PASS, full pkg/... suite 122 ok / 0 FAIL, go build/go vet clean, codegraph bank through the fixed runner 2/2 PASSED real durations + failing case 0/1 FAIL; helix_qa commit 6b46df0). Round 401 added HXC-012 (data race in

helix_code/internal/llm/load_balancer.go background stat-collector goroutine resolved via reproduce-before-fix -race test — Rule 7 / §11.4.43 TDD-fix; SelectOptimalProvider read lb.currentStrategy / lb.strategies on the hot path without holding lb.mutex while SetStrategy writes under the mutex and the background collectStats goroutine mutates lb.stats; fix snapshots the strategy under lb.mutex.RLock() then releases before invoking the strategy and updateStats, observable behaviour unchanged; new internal/llm/load_balancer_race_test.go FAILED RED then PASSES GREEN, full-package go test -race ./internal/llm/... PASS, commit 9d8c1cdc). Round 400 added HXC-006 (HelixCode Speed Programme — all 6 phases / 31 tasks landed making HelixCode 3-5× faster than competitor AI CLI agents; CONST-048 coverage ledger docs/research/speed/05-coverage-ledger.md = 29 PASS + 2 honestly-bounded PARTIAL [P5-T01 build/test parallelism tuning landed but suite-wall-time delta not captured as a pasted benchmark; P5-T02 partial single-provider internal/llm split — Cerebras extracted, full 18-provider split infeasible without an import cycle] + 0 DEFERRED; headline measured wins P1-T02 lazy Ollama 67µs→2.7ns / P1-T03 lazy CLI startup ~8.85× / P2-T03 repo-map cache ~10.6× warm / P2-T06 incremental tree-sitter ~21× / P3-T03 fast-apply ~516× / P4-T01 PGO -46% CPU-bound, every number carrying pasted in-session captured evidence per CONST-035; release-gate sweep — anti-bluff smoke clean, CONST-046 audit no new hardcoded content, governance-cascade verifier's 2 failures are the pre-existing HXC-008 drift not speed regressions; two pre-existing defects surfaced filed honestly as HXC-011 + HXC-012). Round 396 added HXV-003 (LLMsVerifier ProviderAdapterForBenchmark.Complete CONST-050(A) production mock-bluff resolved — WIRE-not-delete decision: the adapter is live code wired by BenchmarkSystem.Initialize, so honest deletion was not applicable; the provider field was retyped from an unused interface{} to the real benchmark.LLMProvider interface, Complete now dispatches to a.provider.Complete returning genuine response text / token count / error, and an honest ErrProviderAdapterNotWired sentinel is returned when no real provider is wired — the hardcoded return "Response", 50, nil BLUFF-001 fabrication is removed; reproduce-before-fix httptest real-dispatch test asserts the adapter actually hits an OpenAI-shim server and returns the real payload; go test ./internal/benchmark/... PASS, go build ./... clean). Round 397 added OPS-001 (LLMOps 2 pre-existing CreatePromptExperiment test failures resolved — both classified (A) test-assertion drift: ControlPromptCreateFails/TreatmentPromptCreateFails asserted pre-idempotency behaviour, expecting a same-(name,version) duplicate to fail CreatePromptExperiment, but commit bb53c38 deliberately made prompt registration idempotent so duplicates are tolerated as no-ops; no production regression; tests re-keyed to the genuine non-tolerated failure path — structurally-invalid prompt with empty Content → registry "prompt content is required" → wrapped with "control prompt"/"treatment prompt" prefix; go test -count=1 ./llmops/... 2 FAIL → 0 FAIL, go build ./...

clean). Round 348 added HXV-002 (LLMsVerifier verification/ package 10 pre-existing test failures resolved — all 10 classified (A) test-assertion drift: 8 verification_test.go Verify tests re-keyed to the round-17 honest ErrVerificationNotWired contract, 2 code_verification_test.go tests re-keyed to honest API-error-propagation / zero-response→failed contracts; no production code changed; go test ./verification/... 10 FAIL → 0 FAIL, go build ./... clean). Round 344 added HXQ-002 (helix_qa pkg/ autonomous ↔ VisionEngine remote API drift resolved — mechanical 3-signature drift: ProbeHosts/SelectStrongestModel/PlanDistribution updated to the round-340 VEN-002 merged superset API, now error-returning with []SSHConfig/[]ModelSpec parameters; 5 config-injected pipeline fields added per CONST-045/046; pkg/autonomous build+test PASS, helix_agent tests/integration compiles). Round 342 added HXA-002 (helix_agent debate/llmprovider sibling-submodule API drift resolved — investigation per operator’s explicit ask found the learning/knowledge/recommendations capability tier was GENUINELY DELETED, not moved: digital.vasic.debate was rebuilt from scratch at commit 196d0ea and the slim CreateDebate/GetStatistics API is the first+only version, with zero surviving copies anywhere in dependencies/; Part-1 import swap + Part-2 slim-API test rewrite + score-scale fix; debate_integration tests PASS, anti-bluff per CONST-035). Round 340 added VEN-002 (VisionEngine vasic-digital-github divergent fork lineage resolved via CONST-061 §11.4.41 merge-first — real 2-parent merge commit 70c9e0c, NO force-push, 16-file conflict surface resolved preserving the anti-bluff truth per CONST-035; 7 packages build+test PASS; 4 remotes fast-forwarded). Round 325 added HXQ-001 (helix_qa TestPerformance flake resolved — 3 pkg/vision/ perf tests gated behind HOST_LOAD_DEDICATED env var; resolution path (b), timing tolerance NOT loosened, anti-bluff strictness preserved per CONST-035).

Coverage by round-system phase

Phase	Closed items	Status
Round 37-89 (governance + stabilization + LLM wiring)	6 batched closures	Done
CONST-046 Phase 1 (rounds 91-93)	3 (pkg/i18n core + audit script + injection wiring)	Done
CONST-046 Phase 2 (rounds 94-96)	3 (SelfImprove + HelixLLM + harmony_os, 15 strings)	Done
	4 closed (DocProcessor +	DONE

Phase	Closed items	Status
CONST-046 Phase 3 (rounds 97-99)	Planning + VisionEngine + panoptic + audit- gate; 20 strings + audit-gate)	
CONST-046 Phase 4 (rounds 100+)	concluded — campaign closed round 463 (HXC-003 Implemented (→ Fixed.md)); genuine user-facing (C) surface exhausted across all 7 scope areas; ~91-462 rounds, tens of thousands of literals	Done

Open issues snapshot (cross-ref docs/ Issues_Summary.md)

2 open issues: 0 Bug / 1 Task (HXC-001) + 1 Operator-blocked Task (HXC-010) / 0 Feature; 0 BLOCKED.

Last regenerated: 2026-05-20 (round 463 — HXC-003 closed Implemented (→ Fixed.md): the CONST-046 i18n migration campaign is concluded — the genuine user-facing (C) string-literal surface is exhausted across all 7 scope areas (helix_code internal/+cmd/+applications/, LLMsVerifier, helix_qa, all owned vasic-digital/+HelixDevelopment/* submodules); ~91-462 rounds migrated tens of thousands of literals through i18n seams with paired-mutation anti-bluff tests; the remaining ~55k audit-baseline hits are all out of CONST-046 scope per docs/audits/2026-05-20-internal-const046-classification.md. Previous round 403 — HXC-008 closed Fixed (→ Fixed.md) + HXC-007/HXC-009 closed Completed (→ Fixed.md). See docs/Fixed.md for full closure entries with commit SHAs.*