

README

Retro-capture: CLI real LLM generate + models list + health

Retro-captured 2026-06-16 vs HEAD 9e6e0458
(9e6e0458f93b7e50ee6d47d8324b34f02e77b31c). This is a CURRENT, fresh end-to-end run of the shipped feature on the running System — it is explicitly a **retro-capture done on 2026-06-16**, NOT a reconstruction of the original feature-shipping moment, and NOT claimed as historical evidence (§11.4.6 / §11.4.123). The original LLM-generate feature is the anti-BLUFF-001/002 resolution in `helix_code/cmd/cli/main.go`; this directory backfills the §11.4.83 docs/qa gap (G7) with a real run.

What was exercised (real end-user CLI path)

Provider: DeepSeek (. ~/api_keys.sh; export
HELIX_LLM_PROVIDER=deepseek).

File	Feature	Real result
cli_generate_run.txt	cli -non-interactive -provider deepseek -prompt "What is 2+2?..."	Model returned 4; real token accounting in=17 out=30 total=47, time=1.298s, finish=length.
cli_models_list.txt	cli -list-models -provider deepseek (CONST-036 LLMsVerifier source-of-truth; anti-BLUFF-002)	Live /models catalog refresh → 2 real models (deepseek-v4-flash, deepseek-v4-pro), real context size 128000. NOT a hardcoded list.
cli_health.txt	cli -health	Real health check: ⚠️ Worker Pool: No healthy workers + ✅ System is operational (honest live state).

Anti-bluff notes

- Every transcript is a REAL run with REAL provider responses — no simulated / placeholder output (§11.4 / §11.4.1 / §11.4.2).
- The `-list-models` run proves models come from a live provider query (catalog refreshed ... (live /models)), not a static array — the canonical BLUFF-002 anti-pattern guard.
- Raft/consensus log noise is the binary's normal single-node startup; the feature output lines (model reply, token table, model list, health verdict) are the load-bearing evidence.