

CLOSE-OUT-CAVEATS

CodeGraph incorporation — Phase C close-out: caveat resolution

Field	Value
Revision	1
Created	2026-05-20
Last modified	2026-05-20
Status	active

Verbatim 2026-05-19 operator mandate: *“all existing tests and Challenges do work in anti-bluff manner - they MUST confirm that all tested codebase really works as expected! We had been in position that all tests do execute with success and all Challenges as well, but in reality the most of the features does not work and can’t be used! This MUST NOT be the case and execution of tests and Challenges MUST guarantee the quality, the completion and full usability by end users of the product!”*

This document closes the three honest caveats left open after the CodeGraph (@colbymchenry/codegraph v0.7.12) incorporation into HelixCode. Authority: HelixCode root CLAUDE.md / CONSTITUTION.md — CONST-035, CONST-050(B), §11.4.3 topology-aware dispatch, §11.4.21 operator-block details.

Table of contents

- [Caveat 1 — Kimi CLI + Qwen Code true end-to-end](#)
- [Caveat 2 — helix_qa autonomous QA session](#)
- [Caveat 3 — anti-bluff covenant in QWEN.md across submodules](#)

Caveat 1 — Kimi CLI + Qwen Code true end-to-end

Both per-agent Challenges were re-driven non-interactively this session against the real CodeGraph MCP server (624,092-node graph of the HelixCode repo). Result: **both agents are genuinely backend-blocked** — reported HONESTLY, not faked.

Kimi CLI — connect-only + tool-discovery PASS (end-to-end operator-blocked)

tools/codegraph/challenges/cg-challenge-05-kimi.sh re-run. The transcript (CG-CHALLENGE-05-kimi-transcript.txt, 1935 bytes) proves Kimi’s MCP loader connected to the codegraph server and enumerated **all 9** codegraph_* tools:

```
MCPServerSnapshot(name='codegraph', status='connected', tools=(  
    'codegraph_search', 'codegraph_context', 'codegraph_callers',
```

```
'codegraph_callees', 'codegraph_impact', 'codegraph_node',  
'codegraph_explore', 'codegraph_status', 'codegraph_files'))
```

The agent step then aborts:

```
Error code: 429 - {'error': {'message': "You've reached kimi monthly usage  
for this billing cycle. ..."}}}
```

This is the strongest available proof short of end-to-end: the MCP transport is fully functional and the tools are reachable; only the LLM backend quota blocks the final answer. Per §11.4.3 the connect-only fallback is reported as connect-only — NEVER as end-to-end.

Qwen Code — connect-only PASS (end-to-end operator-blocked)

tools/codegraph/challenges/cg-challenge-07-qwen.sh re-run. Transcript (CG-CHALLENGE-07-qwen-transcript.txt, 136 bytes):

Qwen OAuth free tier was discontinued on 2026-04-15. Run /auth to switch to Coding Plan, OpenRouter, Fireworks AI, or another provider.

Connect-only PASS captured (config registers codegraph MCP server + binary reachable). End-to-end NOT proven.

Honest disposition

End-to-end Kimi/Qwen verification is **operator-blocked**. Filed as docs/Issues.md → **HXC-010** with ****Status:** Operator-blocked + **Operator-Block-Details:**** per constitution §11.4.21. UNBLOCK CONDITION: operator supplies Kimi API quota/credentials and/or Qwen Code LLM credentials, then re-runs the two Challenges. Claude Code, OpenCode, and Crush remain true-end-to-end verified (Phase C, unchanged).

Caveat 2 — helix_qa autonomous QA session

The codegraph Challenge bank (tools/codegraph/challenges/codegraph-integration.bank.yaml) was driven through the helix_qa runner this session.

- helix_qa/helixqa list -banks <bank> → loads all 7 test cases correctly.
- helix_qa/helixqa run -banks <bank> -platform desktop → the runner loads the 7 definitions (“Auto-wired runner with 7 bridged definitions”) but completes in ~700 ms and reports 0/7 passed with every case **SKIPPED** (see helixqa-session/qa-results/qa-report.md). The runner’s run path is built for Android/UI-driven dumphsys-style platforms and does NOT shell out to a bank’s action: command on the desktop platform — it skips them and emits a hollow “Step Validation PASSED” metadata row (sub-microsecond durations). That hollow PASS is itself a §11.4 PASS-bluff pattern and is NOT counted as evidence here.

Honest disposition. The real, anti-bluff autonomous-session evidence for the codegraph bank is the **direct execution of every bank action: Challenge script**, which each run real codegraph commands / real JSON-RPC / real agent invocations and FAIL LOUDLY on empty/zero/simulated results. Phase C captured that run: tools/codegraph/challenges/run-all.sh → 7/7 PASS with per-challenge wire evidence under this directory (cg-challenge-01..07-*.log, cg-challenge-02-jsonrpc-wire.jsonl, transcripts, connect-proofs). The bank YAML is registered and consumable by the helix_qa runner; what is

NOT yet available is a `helix_qa` run-path that executes `shell-action` banks on the desktop platform. That runner-side gap is a `helix_qa` enhancement, tracked separately — it does not weaken the codegraph anti-bluff evidence, which stands on the direct Challenge-script run with captured wire transcripts.

Captured this session: `helixqa-session/helixqa-run.log + helixqa-session/qa-results/qa-report.md` (the honest record of the runner’s desktop-platform skip behaviour).

Caveat 3 — anti-bluff covenant in QWEN.md across submodules

Assessment of `QWEN.md` presence + anti-bluff covenant across all owned submodules (run 2026-05-20):

Scope	QWEN.md present	CLAUDE.md present
<code>helix_code</code>	no	yes
<code>helix_qa</code>	no	yes
<code>challenges</code>	no	yes
<code>containers</code>	no	yes
<code>security</code>	no	yes
<code>dependencies/ HelixDevelopment/* (10 dirs)</code>	1 of 10 (LLMsVerifier)	10 of 10
<code>dependencies/vasic- digital/* (52 dirs)</code>	0 of 52	52 of 52
Total dependency submodules	1 of 62	62 of 62

Finding: **QWEN.md is effectively absent fleet-wide.** Only one of 67 owned submodules (`submodules/llms_verifier`) ships a `QWEN.md`; every owned submodule already carries a full `CLAUDE.md` with the anti-bluff covenant + `CONST-035` / §11.4 anchors.

Recommendation (NOT executed unilaterally — would mass-create 66+ new files): `QWEN.md` should be made a required consumer governance file fleet-wide ONLY if the **constitution submodule** designates it so (`CONST-059` governs which files are required consumer files — currently the canonical trio is `CONSTITUTION.md` / `CLAUDE.md` / `AGENTS.md`, with `QWEN.md` / `CRUSH.md` named as sibling agent manuals at the meta-repo root by the project `CLAUDE.md` §1.1 but not mandated per-submodule). The single existing submodule `QWEN.md` (`LLMsVerifier`) was inspected — it already carries the anti-bluff covenant mirrored from its `CLAUDE.md`, so no covenant fix was required there.

Disposition: no `QWEN.md` covenant fix needed (the one existing file is already compliant); fleet-wide `QWEN.md` creation deferred pending a constitution-submodule decision on whether `QWEN.md` is a required per-submodule consumer file.