

04-phased-implementation-plan

R4 — HelixCode Speed Programme: Phased Implementation Plan

Revision	1
Created	2026-05-20
Last modified	2026-05-20T00:00:00Z
Status	active
Authority	docs/research/speed/ — operator speed mandate 2026-05-20

Table of contents

- [1. Scope, conventions & how to read a task](#)
- [2. Phase-dependency graph](#)
- [3. Phase 0 — Measurement baseline](#)
- [4. Phase 1 — Low-risk high-impact LLM & startup wins](#)
- [5. Phase 2 — Context-build speed](#)
- [6. Phase 3 — Interactive & agent-loop levers](#)
- [7. Phase 4 — Profile-gated tuning & caching architecture](#)
- [8. Phase 5 — Developer-experience & submodule cascade](#)
- [9. §11.4.74 submodule-catalogue check \(mandatory before any new component\)](#)
- [10. Per-task test-floor & anti-bluff template](#)
- [11. Task summary table](#)

1. Scope, conventions & how to read a task

This plan executes the unified opportunity list O1–O21 from [00-speed-programme-overview.md](#) §5. It divides the work into **6 phases (Phase 0 through Phase 5)** containing **31 tasks total**.

One task ≈ one subagent-sized work unit (≈1 PWU per constitution §11.4.58) — small enough that a single subagent-driven round can implement, test, prove and commit it. Tasks are ordered within a phase by dependency then impact-to-effort.

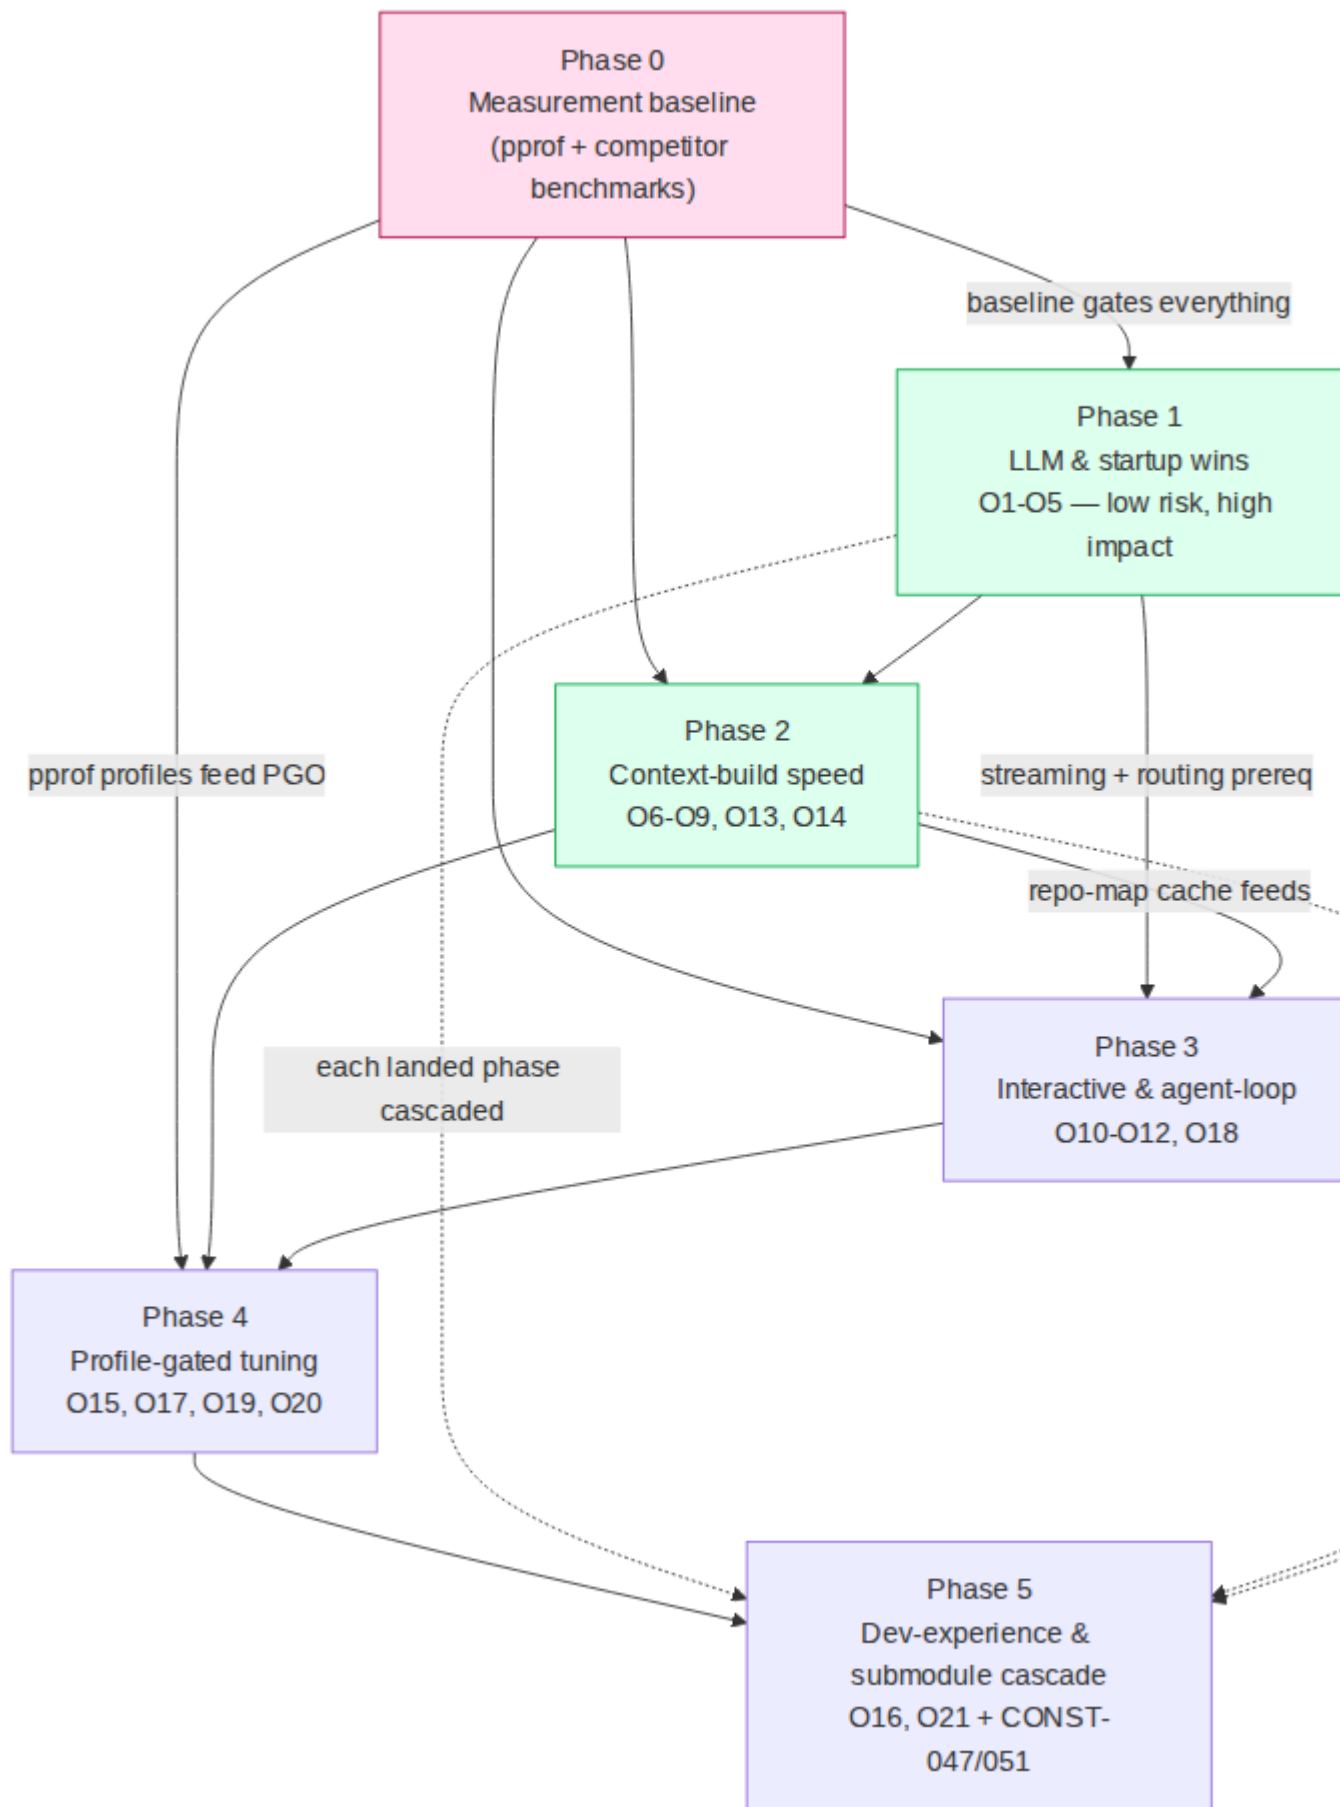
Each task is specified as:

- **ID** — Pn-Tnn (phase, task).
- **Title.**
- **Target files** — the real `helix_code/...` paths to change (from R1 evidence).
- **Implements** — cross-reference to R1 bottleneck / R2 technique / R3 technique / R4 opportunity (O-number).
- **Expected speedup** — the claim the benchmark must prove.
- **Risk** — Low / Med / High (regression risk to working features).
- **Test types** — per CONST-050: unit + integration + benchmark + Challenge + HelixQA. (Mocks permitted ONLY in unit tests — CONST-050(A).)
- **Anti-bluff proof** — the captured runtime evidence each test must produce (CONST-035 / Article XI §11.9). “Faster” without pasted before/after numbers is a bluff (Rule 9).

Path note (R1 §2): `helix_code/` IS the inner Go module (module `dev.helix.code`, go 1.26). The CLAUDE.md text describing a deeper `helix_code/helix_code/` path is stale — paths below use the confirmed layout.

Hard constraint: no working feature may regress (R4 §6). Every task carries the no-regression proof — the functional tests that passed before still pass after.

2. Phase-dependency graph



Mermaid diagram 1

Critical path: Phase 0 → Phase 1 → Phase 2 → Phase 4 → Phase 5. Phase 0 is a hard gate — no speedup claim is falsifiable before the baseline exists (R4 §7). Phase 3 depends on Phase 1 (streaming + provider plumbing) and Phase 2 (warm repo-map cache feeds context for the agent loop). Phase 5 cascades each landed phase's work into owned submodules and may begin its per-submodule cascade as soon as the corresponding meta-repo phase is green (dotted edges).

3. Phase 0 — Measurement baseline

Goal: capture the before-state. Nothing else in the programme can be validated until this phase produces committed baseline artefacts. No production code changes in this phase — it builds the measurement harness and captures numbers.

P0-T01 — pprof harness for the four hot paths

- **Target files:** `helix_code/cmd/cli/main.go` (add opt-in `--pprof` flag wiring), `helix_code/internal/server/` (verify `net/http/pprof` mount), new `helix_code/tests/performance/pprof_harness_test.go`.
- **Implements:** R3 §1.2; R1 §6 items 1–4; R4 §7.
- **Expected speedup:** none (enabling task).
- **Risk:** Low.
- **Test types:** unit (harness wiring); integration (pprof endpoint serves a profile); benchmark (the harness runs the 4 scenarios); Challenge (capture-a-profile workflow); HelixQA (perf bank registers the harness).
- **Anti-bluff proof:** committed `.pprof` files for S1–S4 under `docs/research/speed/baseline/`; pasted `go tool pprof -top` output.

P0-T02 — Benchmark suite for startup / llm-dispatch / repomap / search

- **Target files:** new `*_test.go` benchmarks in `helix_code/cmd/cli/`, `helix_code/internal/llm/`, `helix_code/internal/repomap/`, `helix_code/internal/tools/filesystem/`.
- **Implements:** R1 §6 item 3 (these hot paths currently lack benchmarks); R3 §8.3.
- **Expected speedup:** none (enabling).
- **Risk:** Low.
- **Test types:** benchmark (the suite itself); unit (benchmark helpers); HelixQA (benchmark bank registration).
- **Anti-bluff proof:** pasted `go test -bench -benchmem` output for each new benchmark, committed as `baseline/benchmarks-2026-05-20.txt`.

P0-T03 — Competitor wall-clock baseline on identical hardware

- **Target files:** new scripts/testing/competitor_speed_baseline.sh; results under docs/research/speed/baseline/.
- **Implements:** R1 §6 item 6; R2 §2 (Go-rewrite 68× data point).
- **Expected speedup:** none (enabling — sets the number to beat).
- **Risk:** Low.
- **Test types:** integration (script runs each agent for S1–S4); Challenge (reproducible competitor-benchmark run); HelixQA (bank).
- **Anti-bluff proof:** committed table of measured wall-clock for Claude Code / Gemini CLI / Aider / Cline / Crush across S1–S4; raw /usr/bin/time output.

P0-T04 — Scenario fixtures + the canonical-scenario runner

- **Target files:** new helix_code/tests/performance/scenarios/ (large-repo fixture generator), scripts/testing/run_speed_scenarios.sh.
- **Implements:** R4 §7 (S1–S4 canonical scenarios).
- **Expected speedup:** none (enabling).
- **Risk:** Low.
- **Test types:** integration (runner produces stable numbers across 3 runs); Challenge (end-to-end scenario run); HelixQA (bank).
- **Anti-bluff proof:** pasted 3-run variance for each scenario proving the harness is stable enough to detect a 1.3×+ change.

4. Phase 1 — Low-risk high-impact LLM & startup wins

Goal: the S-tier wins (R3 sequencing Wave 1) — ~2–3× perceived-latency from prompt caching + HTTP tuning + streaming + lazy startup. Low regression risk.

P1-T01 — Shared tuned HTTP/2 transport across all providers

- **Target files:** new helix_code/internal/providers/httpclient.go (shared factory); edit helix_code/internal/llm/{openai,deepseek,gemini,azure,copilot,cerebras,koboldai,local}_provider.go + local_llm_manager.go.
- **Implements:** R1 B03; R2 #9; R3 §4.7; O2.
- **Expected speedup:** removes per-call TLS handshake on bursts; high on concurrent/rapid-fire requests.
- **Risk:** Low.
- **Test types:** unit (factory returns tuned transport — MaxIdleConnsPerHost≥32, ForceAttemptHTTP2); integration (real provider call reuses a connection); benchmark (burst of N calls, connection-

reuse count); Challenge (real LLM generate over a warm pool); HelixQA (provider bank).

- **Anti-bluff proof:** captured netstat/transport idle-conn count before/after proving reuse; pasted benchmark delta.

P1-T02 — Lazy/async Ollama discovery out of the constructor

- **Target files:** `helix_code/internal/llm/ollama_provider.go` (remove `discoverModels()` from `NewOllamaProvider`, line ~109), `helix_code/cmd/cli/main.go` (~234).
- **Implements:** R1 B02; R2 #6; R3 §3.1; O5.
- **Expected speedup:** removes one synchronous HTTP round-trip from every CLI start.
- **Risk:** Low.
- **Test types:** unit (constructor does no network I/O — mock asserts zero calls); integration (discovery still works on first `GetModels`); benchmark (`NewCLI()` wall-clock before/after); Challenge (`--list-models` still returns real models); HelixQA (bank).
- **Anti-bluff proof:** pasted `NewCLI()` benchmark delta; Challenge output showing real Ollama models still listed (no regression).

P1-T03 — Lazy CLI startup: sync.Once getters on *CLI

- **Target files:** `helix_code/cmd/cli/main.go` `Run()` (lines ~459–680) — convert ~20 eager subsystems to `sync.Once`-guarded lazy getters.
- **Implements:** R1 B01, B13, B14, B18; R2 #6; R3 §3.1, §3.2; O4.
- **Expected speedup:** high cold-start — short commands skip worktree git shell-out, hooks YAML, LSP LookPath sweep, MCP spawn.
- **Risk:** Med (touches the bootstrap of every command — needs the full no-regression suite).
- **Test types:** unit (each getter constructs once, lazily); integration (every command path still wires its subsystems on demand); benchmark (S1 cold start); Challenge (each top-level command still works end-to-end); HelixQA (CLI bank, full command matrix).
- **Anti-bluff proof:** pasted S1 wall-clock before/after; the full pre-change CLI command test matrix passing post-change (no-regression proof).

P1-T04 — Prompt-cache-stable prefix + deterministic JSON serialization

- **Target files:** `helix_code/internal/llm/request-builders` (system prompt + tool-definition assembly), new `helix_code/internal/llm/promptcache/` (cache-break detection + deterministic key ordering).

- **Implements:** R2 #1, #8; R3 §4.1; O1. Go pitfall — map-key order randomization must be eliminated.
- **Expected speedup:** up to ~85% TTFT on long prompts (cache hit).
- **Risk:** Low-Med (prefix changes must be invisible to feature behaviour).
- **Test types:** unit (serialization is byte-deterministic across 1000 runs; cache-break detector flags a mutated prefix); integration (real Anthropic call reports `cache_read_input_tokens > 0` on the 2nd request); benchmark (TTFT cold vs warm); Challenge (multi-turn session, 2nd turn hits cache); HelixQA (bank).
- **Anti-bluff proof:** pasted API response showing `cache_creation_input_tokens` then `cache_read_input_tokens` on turn 2; TTFT delta.

P1-T05 — `cache_control` opt-in for every supported provider

- **Target files:** `helix_code/internal/llm/{anthropic,openai,gemini,...}_provider.go` request paths.
- **Implements:** R2 #8 (Plandex all-provider); R3 §4.1; O1.
- **Expected speedup:** extends T04's win to OpenAI/Gemini (implicit + explicit caching).
- **Risk:** Low-Med.
- **Test types:** unit (each provider emits provider-correct cache directives); integration (real call per provider reports cache hit); benchmark (per-provider TTFT cold/warm); Challenge; HelixQA (multi-provider bank).
- **Anti-bluff proof:** per-provider pasted cache-hit metric from the real API.

P1-T06 — Cache pre-warming at session start

- **Target files:** `helix_code/internal/session/` (session-open hook), `helix_code/internal/llm/promptcache/`.
- **Implements:** R3 §4.2; O1.
- **Expected speedup:** removes the first-request cold-cache penalty.
- **Risk:** Low.
- **Test types:** unit (warm request is `max_tokens-minimal`); integration (first real user request after warm-up is a cache hit); benchmark (first-turn TTFT with vs without pre-warm); Challenge; HelixQA (bank).
- **Anti-bluff proof:** pasted first-turn `cache_read_input_tokens > 0`.

P1-T07 — Streaming-first verification on every surface

- **Target files:** `helix_code/cmd/cli/main.go`, `helix_code/applications/terminal-ui/`, `helix_code/applications/desktop/` — verify all consume `GenerateStream`.
- **Implements:** R2 #9; R3 §4.3; O3. (Verify, do not regress — mostly present.)

- **Expected speedup:** high perceived (time-to-first-visible-token).
- **Risk:** Low.
- **Test types:** unit (each surface renders incrementally); integration (real streamed response renders token-by-token); benchmark (time-to-first-token); Challenge (CLI + TUI + desktop each show streaming); HelixQA (UI bank).
- **Anti-bluff proof:** captured timestamped render log proving first token shown before completion arrives.

5. Phase 2 — Context-build speed

Goal: R3 Wave 2 context wins — $\sim 1.5\text{--}2\times$ on context-build phases via repo-map cache + parallelism + incremental parsing.

P2-T01 — `filepath.WalkDir` migration + buffered I/O

- **Target files:** the 19 `filepath.Walk` sites (R1 B09) — `internal/repomap/repomap.go:248`, `internal/context`, `internal/fix`, `internal/rules`, `internal/tools/{mapping,web,filesystem}`, `internal/workflow`, etc.
- **Implements:** R1 B09; R3 §7.1, §7.2; O13.
- **Expected speedup:** medium on large trees (cheaper `fs.DirEntry`).
- **Risk:** Low (mechanical).
- **Test types:** unit (walk yields identical file set); integration (repo-map + search still cover the same files); benchmark (cold walk of S2 fixture); Challenge; HelixQA (bank).
- **Anti-bluff proof:** pasted before/after walk benchmark; file-set equality test.

P2-T02 — Hoist per-call regexp; remove redundant body copies

- **Target files:** `internal/editor/formats/{editor,diff,architect,whole}_format.go` (≥ 17 `MustCompile` sites — R1 B12); `internal/llm/ollama_provider.go:359,398` (`bytes.NewReader` not `strings.NewReader(string(...))` — R1 B11).
- **Implements:** R1 B11, B12; R3 §1.5; O14.
- **Expected speedup:** medium on edit-format parse hot path.
- **Risk:** Low.
- **Test types:** unit (package-level regexes compile once; parse output unchanged); integration (real LLM-edit-format parse); benchmark (parse loop alloc/op); Challenge; HelixQA (bank).
- **Anti-bluff proof:** pasted -benchmem allocs/op before/after.

P2-T03 — Persistent content-addressed repo-map cache

- **Target files:** `helix_code/internal/repomap/cache.go` (key by path+mtime/ content-hash; one background writer replacing per-Set goroutine — R1 B07), `repomap.go:148` (stat-based token estimate — R1 B19).
- **Implements:** R1 B07, B15, B19; R2 #2 (Aider diskcache + getmtime); R3 §5.2, §5.4, §6.1; O6.
- **Expected speedup:** ~1.5–2× on warm context (unchanged files never re-parsed).
- **Risk:** Med (cache correctness — stale entry must miss).
- **Test types:** unit (changed file → cache miss; unchanged → hit); integration (warm second index reuses cache); benchmark (cold vs warm S2 index); Challenge (edit one file, re-index, only that file re-parsed); HelixQA (bank).
- **Anti-bluff proof:** pasted cold-vs-warm index wall-clock; instrumentation log showing N-1 cache hits after a 1-file edit.

P2-T04 — Parallelise repo-map: worker pool + parser pool + single-pass stats

- **Target files:** `helix_code/internal/repomap/repomap.go` (consume the unused `MaxConcurrency`, line 29 — R1 B04; single-pass `GetStatistics` — R1 B05; narrow lock — R1 B16), `tree_sitter.go:56` (`sync.Pool` of parsers — R1 B06).
- **Implements:** R1 B04, B05, B06, B16; R2 #2; R3 §2.1, §2.2, §1.4; O7.
- **Expected speedup:** high on cold index (tree-sitter is the CPU bottleneck).
- **Risk:** Med (concurrency correctness — race-tested).
- **Test types:** unit (-race; worker pool bounded by `MaxConcurrency`; parser pool resets between files); integration (parallel index = serial index output); benchmark (cold S2 index, varying concurrency); Challenge; HelixQA (bank).
- **Anti-bluff proof:** pasted -race clean run; output-equality test (parallel vs serial); benchmark delta vs P0 baseline.

P2-T05 — Parallelise the SearchContent grep tool

- **Target files:** `helix_code/internal/tools/filesystem/searcher.go:243` — bounded worker pool (`errgroup.SetLimit`) over candidate paths.
- **Implements:** R1 B08; R2 #2 (Continue ripgrep); R3 §2.1, §7.2; O8.
- **Expected speedup:** high on large trees (the agent Grep hot path).
- **Risk:** Low-Med (result ordering — must remain deterministic).
- **Test types:** unit (parallel search = serial search result set, sorted); integration (real grep over S3 fixture); benchmark (S3 wall-clock); Challenge (agent Grep tool call end-to-end); HelixQA (tools bank).
- **Anti-bluff proof:** result-set equality test; pasted S3 benchmark delta.

P2-T06 — Incremental tree-sitter parsing (edit API)

- **Target files:** `helix_code/internal/repomap/tree_sitter.go` — use the `smacker/go-tree-sitter` incremental edit API instead of full re-parse.
- **Implements:** R1 B06 (extends); R2 #2; R3 §5.1, §5.3; O9.
- **Expected speedup:** high on per-edit context refresh (re-parse only edited regions).
- **Risk:** Med (incremental-parse correctness vs full re-parse).
- **Test types:** unit (incremental re-parse == full re-parse AST for an edited file); integration (edit → re-index produces correct symbols); benchmark (per-edit re-parse vs full); Challenge; HelixQA (bank).
- **Anti-bluff proof:** AST-equality test (incremental vs full); pasted per-edit benchmark delta.

P2-T07 — Config loaded once, threaded down

- **Target files:** `helix_code/internal/config/config.go:215 (Load())`, `helix_code/cmd/cli/main.go:241`, server + subagent config paths.
- **Implements:** R1 B10; R3 §3.4; O16.
- **Expected speedup:** low-medium (drops repeat YAML reads / Viper global churn).
- **Risk:** Low-Med (config plumbing — every consumer must get the same struct).
- **Test types:** unit (`Load()` called once; struct threaded); integration (CLI + server + subagent see consistent config); benchmark (startup config-load count); Challenge; HelixQA (bank).
- **Anti-bluff proof:** instrumentation log proving a single `ReadInConfig` per process; no-regression config tests.

6. Phase 3 — Interactive & agent-loop levers

Goal: the interactive multipliers — model routing, fast-apply, tool parallelism, compaction. Depends on Phase 1 (provider plumbing) + Phase 2 (repo-map context).

P3-T01 — Small-model routing for cheap subtasks (cascade)

- **Target files:** new `helix_code/internal/llm/routing/`; callers in `internal/agent/`, `internal/commands/` (classification, ranking, commit-msg, ambiguity detection).
- **Implements:** R2 #7; R3 §4.4; O10. Model metadata MUST come from `LLMsVerifier (CONST-036/037)` — no hardcoded model lists.
- **Expected speedup:** high on multi-step agent loops (cheap tasks on fast models).

- **Risk:** Med (routing must not degrade output quality — escalate on low confidence).
- **Test types:** unit (routing policy picks small model for trivial tasks; mock verifier); integration (real cascade — small model first, escalate); benchmark (agent-loop wall-clock with vs without routing); Challenge (real multi-step task); HelixQA (bank).
- **Anti-bluff proof:** captured per-subtask model-used log; loop wall-clock delta; quality-parity assertion vs frontier-only.

P3-T02 — Diff-style edits (SEARCH/REPLACE only changed lines)

- **Target files:** `helix_code/internal/editor/formats/` — emit only changed lines; `internal/tools/` apply path with fuzzy match.
- **Implements:** R2 #3, #4 (Roo `apply_diff`, Aider `diff/udiff`); R3 (edit format); O11.
- **Expected speedup:** ~30% fewer output tokens → proportional latency cut.
- **Risk:** Med (malformed diff = correctness + speed loss — benchmark edit-format compliance).
- **Test types:** unit (diff round-trips; malformed diff rejected); integration (real LLM emits a valid diff that applies); benchmark (output-token count full-rewrite vs diff); Challenge (real multi-edit); HelixQA (edit bank).
- **Anti-bluff proof:** pasted output-token delta; edit-format compliance rate.

P3-T03 — Dedicated fast-apply path

- **Target files:** new `helix_code/internal/editor/fastapply/`; integrates with P3-T01 routing.
- **Implements:** R2 #3 (Cursor speculative edits 9–13×, Morph 4500–10,500 tok/s); R3 §4.5; O11.
- **Expected speedup:** high interactive — file apply ~10× faster.
- **Risk:** Med-High (apply correctness is non-negotiable — full no-regression edit suite).
- **Test types:** unit (fast-apply output == reference apply); integration (real apply-model / speculative path produces a correct file); benchmark (apply tok/s); Challenge (real file-edit workflow); HelixQA (edit bank).
- **Anti-bluff proof:** byte-equality of fast-apply vs reference apply on a corpus; pasted apply tok/s before/after.

P3-T04 — Aggressive tool-call parallelism

- **Target files:** `helix_code/internal/agent/` tool-dispatch loop; `internal/tools/registry.go`.
- **Implements:** R2 #5 (Claude Code /batch up to 10×); R3 §2.1; O12.
- **Expected speedup:** high on multi-file / multi-tool turns.

- **Risk:** Med (independent-tool detection — only parallelise side-effect-free or non-conflicting calls).
- **Test types:** unit (independent calls run concurrently, dependent serialised); integration (real multi-tool turn); benchmark (multi-tool turn wall-clock); Challenge (real multi-file task); HelixQA (agent bank).
- **Anti-bluff proof:** captured concurrency trace; turn wall-clock delta; output-equality vs serial.

P3-T05 — History condenser / compaction

- **Target files:** `helix_code/internal/context/`, `internal/session/` — summarise history to keep the cached prefix small.
- **Implements:** R2 #10 (OpenHands condenser, `/compact`); R4 O18.
- **Expected speedup:** medium on long autonomous runs (prevents context bloat).
- **Risk:** Med (compaction must not drop task-critical context).
- **Test types:** unit (condenser preserves key facts); integration (long session stays under the window); benchmark (long-run turn latency with vs without); Challenge (real long autonomous run); HelixQA (bank).
- **Anti-bluff proof:** captured token-count trajectory; task-success parity long-run with vs without compaction.

7. Phase 4 — Profile-gated tuning & caching architecture

Goal: R3 Wave 3 — only where Phase 0 pprof evidence proves the bottleneck (§11.4.6 no-guessing).

P4-T01 — PGO with a real production default.pgo

- **Target files:** `helix_code/cmd/cli/default.pgo` (committed), `helix_code/cmd/server/default.pgo`, `helix_code/Makefile` (PGO build note).
- **Implements:** R3 §1.1; O15. Profile sourced from Phase 0 P0-T01.
- **Expected speedup:** 2–14% CPU-bound.
- **Risk:** Low.
- **Test types:** unit (build picks up `.pgo`); benchmark (S2/S3 with vs without PGO); integration (PGO binary passes the full suite); Challenge; HelixQA (bank).
- **Anti-bluff proof:** pasted benchmark delta PGO vs non-PGO; committed profile.

P4-T02 — Multi-tier cache (memory → disk → Redis) for context/embeddings

- **Target files:** `new helix_code/internal/cache/` (L1/L2/L3); integrate with `internal/repomap/cache.go`, `internal/redis/`, embedding paths.
- **Implements:** R2 #2; R3 §6.1, §5.6; O17. §11.4.74 check — see §9.
- **Expected speedup:** high (warm reads ~0.1 ms L1).
- **Risk:** Med (cache coherence across tiers).
- **Test types:** unit (L1 → L2 → L3 read-through; populate-upward); integration (real Redis L3); benchmark (read latency per tier); Challenge (warm-cache workflow); HelixQA (bank).
- **Anti-bluff proof:** per-tier latency measurement; coherence test (write invalidates all tiers).

P4-T03 — Config-driven DB pool sizing

- **Target files:** `helix_code/internal/database/database.go:45-48` (pool MaxConns/MinConns from config; smaller CLI default).
- **Implements:** R1 B17; O19.
- **Expected speedup:** low (fewer idle conns for CLI).
- **Risk:** Low.
- **Test types:** unit (pool sizing reads config); integration (real PostgreSQL, pool honours config); benchmark (CLI startup conn count); Challenge; HelixQA.
- **Anti-bluff proof:** pasted pool-stat output proving config-driven sizing.

P4-T04 — Profile-gated alloc / GC / contention tuning

- **Target files:** profile-confirmed hot paths only — candidates `internal/memory/memory_manager.go` (32 locks — R1 B22), `internal/cognee/service.go` (R1 B23); `sync.Pool` / sharded locks / `RWMutex` per pprof.
- **Implements:** R1 B22, B23; R3 §1.3, §1.4, §1.7, §2.3; O20.
- **Expected speedup:** 1.05–1.3× on CPU/lock-bound slices.
- **Risk:** Med (concurrency correctness — -race mandatory).
- **Test types:** unit (-race; pool/lock correctness); integration (concurrent agent turns); benchmark (mutex-profile contention before/after); Challenge; HelixQA (bank).
- **Anti-bluff proof:** pasted mutex/block profile before/after proving the bottleneck was real and reduced; -race clean.

8. Phase 5 — Developer-experience & submodule cascade

Goal: developer-experience speed + recursive application of every landed speedup to owned submodules (CONST-047, CONST-051 — submodules are equal-codebase).

P5-T01 — Build cache, -p/GOMAXPROCS, test parallelism tuning

- **Target files:** helix_code/Makefile (test targets — R1 B20), Makefile.
- **Implements:** R1 B20; R3 §8.1, §8.2, §8.3; O21.
- **Expected speedup:** developer-experience (suite wall-time).
- **Risk:** Low.
- **Test types:** integration (suite still passes with tuned -p/-parallel); benchmark (suite wall-time before/after); Challenge; HelixQA.
- **Anti-bluff proof:** pasted suite wall-time delta; full suite still green.

P5-T02 — Sub-package internal/llm for incremental-compile speed

- **Target files:** helix_code/internal/llm/ → internal/llm/providers/... (30k LOC / 75 files — R1 B21).
- **Implements:** R1 B21; R3 §8.4; O21.
- **Expected speedup:** developer-experience (smaller recompile unit).
- **Risk:** Med (import-path churn — every importer updated atomically).
- **Test types:** unit (all packages compile + pass); integration (no behaviour change); benchmark (incremental compile after a 1-line edit); Challenge; HelixQA.
- **Anti-bluff proof:** pasted incremental-compile time before/after; full suite green; no import-path drift.

P5-T03 — Cascade Phase 1–4 speedups into owned submodules

- **Target files:** every owned submodule under vasic-digital / HelixDevelopment (and the other authorised orgs per CONST-051) that contributes to a measured scenario — helix_qa/, challenges/, security/, etc.
- **Implements:** CONST-047, CONST-051(A) (equal-codebase); R4 §6.
- **Expected speedup:** programme-wide — each submodule gets the same posture.
- **Risk:** Med (per-submodule regression surface).
- **Test types:** per submodule — unit + integration + benchmark + Challenge + HelixQA; recursive per CONST-047.

- **Anti-bluff proof:** per-submodule before/after benchmark; per-submodule no-regression suite; coverage-ledger entry per CONST-048.

P5-T04 — Speed-programme coverage ledger & release-gate sweep

- **Target files:** docs/research/speed/ ledger; scripts/generate-coverage-ledger.sh extended for speed metrics.
- **Implements:** CONST-048 (coverage ledger — feature × platform × invariant).
- **Expected speedup:** none (governance close-out).
- **Risk:** Low.
- **Test types:** integration (ledger regenerates); Challenge (ledger sweep); HelixQA.
- **Anti-bluff proof:** committed ledger showing every O1–O21 task with captured before/after evidence and a green test floor.

9. §11.4.74 submodule-catalogue check (mandatory before any new component)

Per constitution §11.4.74, **before any task introduces a new reusable component, the vasic-digital / HelixDevelopment owned-submodule catalogue MUST be checked** — discover dynamically via `gh repo list vasic-digital`, `gh repo list HelixDevelopment` and `glab` for the GitLab mirrors. Incorporate an existing owned component rather than rebuilding it (CONST-051(C) — flat layout, no nested own-org chains; CONST-054 — record in `helix-deps.yaml`).

Tasks that introduce a candidate-reusable component and therefore REQUIRE the catalogue check before implementation:

Task	New component	Check before building
P0-T01	pprof profiling harness	helix_qa/ perf tooling; R3 §1.1 note.
P1-T01	shared HTTP-client factory	helix_code/internal/providers; R3 §4.7 note.
P2-T04	bounded worker-pool / scheduler	helix_code/internal/worker; owned-submodule catalogue; R3 §2.2 note.
P2-T03 / P4-T02	content-addressed / multi-tier cache	helix_code/internal/redis; owned-submodule catalogue; R3 §6.1 note.
P3-T01	model-routing component	LLMsVerifier (CONST-036) is the model-metadata source — do not duplicate.
P3-T03	fast-apply component	owned-submodule catalogue; R2 #3.

If the catalogue has no suitable component, the new component is built **decoupled and project-not-aware** (CONST-051(B)) so it can itself become a catalogue entry.

10. Per-task test-floor & anti-bluff template

Every task above is closed only when ALL of the following are captured in the task’s commit/PR (CONST-035, CONST-048, CONST-050):

- 1. **Unit tests** — mocks permitted *only here* (CONST-050(A)). Prove the changed unit behaves; for concurrency tasks, -race clean.
- 2. **Integration tests** — real infrastructure (real PostgreSQL/Redis/Ollama/LLM provider), no mocks. Prove the change works in the real system.
- 3. **Benchmark** — go test -bench -benchmem, **before AND after**, pasted. The delta IS the speedup claim. No delta = bluff (CONST-035 Rule 9).
- 4. **Challenge** — a challenges/ script exercising the real end-to-end user workflow the task touches (CONST-006/050).
- 5. **HelixQA** — the change registered in the relevant helix_qa/ test bank; autonomous QA session run with captured wire evidence (CONST-050).
- 6. **No-regression proof** — the functional tests that passed before the change still pass after it (R4 §6 — no working feature regresses).

Anti-bluff statement each task’s evidence must satisfy: pasted terminal output from *that session* showing (a) the before number, (b) the after number, (c) the ratio meeting the task’s “Expected speedup”, (d) the full functional suite green. Words like *verified / tested / faster / complete* without this output are forbidden (CLAUDE.md Rule 9).

11. Task summary table

Phase	Tasks	Theme	Risk band
Phase 0	P0-T01 ... P0-T04 (4)	Measurement baseline — pprof + benchmarks + competitor wall-clock + scenario harness	Low
Phase 1	P1-T01 ... P1-T07 (7)	LLM & startup wins — HTTP tuning, lazy startup/Ollama,	Low–Med

Phase	Tasks	Theme	Risk band
		prompt cache, streaming	
Phase 2	P2-T01 ... P2-T07 (7)	Context-build speed — WalkDir, regex hoist, repo-map cache + parallelism, incremental parse, config-once	Low–Med
Phase 3	P3-T01 ... P3-T05 (5)	Interactive & agent-loop — model routing, diff edits, fast-apply, tool parallelism, compaction	Med–High
Phase 4	P4-T01 ... P4-T04 (4)	Profile-gated tuning — PGO, multi-tier cache, DB pool, alloc/GC/ contention	Low–Med
Phase 5	P5-T01 ... P5-T04 (4)	Dev-experience & submodule cascade — build/test tuning, sub-packaging, CONST-047/051 cascade, coverage ledger	Low–Med
Total	31 tasks	6 phases	Phase 0 gates all; low-risk-first ordering

Execution: each task is one subagent-driven round per the operator’s standing preference. Phases run in the dependency order of §2; within a phase, tasks run in listed order (dependency then impact-to-effort). Phase 0 must complete and commit its baseline artefacts before any Phase 1+ speedup claim is admissible.