

2026-06-10-impl-D17-debluff-review

Independent Code-Review — D-16 + D-17 De-Bluff of the per-agent CLI layer

Reviewer: Independent code-review subagent (§11.4.125 / §11.4.142) **Date:** 2026-06-10 **Repo:** HelixCode root / Volumes/T7/Projects/HelixCode **Target:** UNCOMMITTED working-tree changes in submodules/helix_agent/internal/clis/agents/
Verdict: GO

1. Scope verification (step 6) — PASS

- 122 changed entries total. **Every** path is under `internal/clis/agents/`.
- `git status --porcelain` confirms: 92 modified `.go` files (61 source + 31 `_test.go`) + 30 untracked test files (`*_pin_test.go` / `*_debluff_pin_test.go` / `*_stub_pin_test.go` + one `copilotworkspace_test.go`). 29 of the untracked files are `pin/debluff/stub` guard tests.
- **NOT touched (correct):** `instance_manager.go`, `go.mod`, `go.sum`, anything in `helix_code/`. Confirmed via full porcelain listing.
- **Secret scan:** `git diff` over the batch — no `sk-...`, AWS keys, inline `api_key=` / `password=`. CLEAN.
- **Mutation-residue scan (§11.4.84):** no MUTATED for `paired///` always `pass/_mutated_` markers in non-test source. CLEAN. (One false-positive `grep` hit, `"Task completed successfully"` in `smolagents.go:199`, is a doc-comment *quoting* the removed bluff — benign.)
- Working tree restored exactly after all mutation cycles: still 122 files, build green.

2. Whole-layer build + vet (step 1) — PASS (the gate)

```
cd submodules/helix_agent
go build ./internal/clis/... → BUILD_EXIT=0      (no output)
go vet   ./internal/clis/... → VET_EXIT=0        (no output)
```

No cross-package breakage from the ~60 parallel edits. The layer compiles and vets together.

3. Whole-layer test (step 2) — PASS (the gate)

```
go test ./internal/clis/... -count=1
```

Result: ALL ok, 0 FAIL. 78 packages reported ok (incl. every touched agent package: `amazonq`, `claude_code`, `cline`, `codai`, `codeiumwindsurf`, `codenamegoose`, `cody`, `conduit`, `continueagent`, `copilotcli`, `copilotworkspace`, `crush`, `cursor`, `deepseek(+cli)`, `fauxpilot`, `forge`, `gemini`, `getshitdone`, `githubspark`, `gitmcp`, `gptengineer`, `gptme`, `gpnr`, `honeycomb`, `hunyuan`, `jetbrainsai`, `junie`, `kilocode`, `kimi`, `kodu`, `lovable`, `mistralcode`, `mobileagent`, `multiagentcoding`, `nanocoder`, `noi`, `octogen`, `ollamacode`, `opencodecli`, `perplexity`, `plandex`, `postgresmcp`, `promptfoo`,

shai, smolagents, snowcli, speckit, supermaven, superset, tabnine, taskweaver, uiuxpromax, vtcode, warp, windsurf). 4 packages report [no test files] (claude/api, claude/features, claude/strategy) — pre-existing, untouched.

4. Anti-bluff invariant (step 3) — PASS

- **CLAUDE.md §9 smoke (scoped to internal/clis/agents):** 2 matches, both benign:
 - `claude_code.go:247` — a `// simulated/templated response (BLUFF-001)` . doc-comment describing what the code **MUST NOT** do.
 - `claude_code/mcp_integration.go:26` — `ErrMCPCClientNotWired` error string quoting the historical bluff. (This file is committed/out-of-scope, not part of this batch.)
 - No **live** fabrication literal in non-test source.
- **Targeted response-fabrication scan** (`fmt.Sprintf("...received|Generated by|Reviewed|Applied edit|completed successfully...")` as a `RESPONSE`): `CLEAN` across agents source.
- **Broad fabricated-status-map scan** (`"status": "generated|sent|executed|composed|..."`): the only hits in agents source are in `agentdeck.go` and `bridle.go` — **both out of scope** (not in this batch's diff; pre-existing).
- **Spot-check of 8 packages confirms genuine conversion:**
 - **Real-exec class** — `claude_code` (`resolveClaudeBinary`→`exec.LookPath`, `runClaude`→`exec.CommandContext+CombinedOutput`, `CLAUDE_BIN` test override; removed `"Claude Code received: %s"/"Applied edit to..."/"Reviewed..."`), `gemini` (`geminiBinary/GEMINI_BIN`, `exec.LookPath+CommandContext+CombinedOutput`; removed `"// Generated by Gemini"` template). Remaining `fmt.Sprintf("Review %s ...") / "Edit file %s: %s"` are **prompts passed INTO the binary**, not fabricated responses — correct.
 - **Honest-error class** — `cursor` (`errNoHeadlessCLI`; preserves genuine session CRUD, removes only the fabricated assistant turn), `cline` (`errNoHeadlessCLI`; `exec.LookPath` retained only for `IsAvailable` detection), `copilotworkspace` (`errNoHeadlessCLI`), `gptengineer` (`ErrCLINotWired`), `smolagents/superset/promptfoo` (`inline fmt.Errorf("... refusing to fabricate ... (BLUFF-001)"`) — removed fabricated transcripts / scorecards / dashboard URLs).
- **Per-package completeness check:** every modified source file carries either a real `exec.*` call or an honest-error return citing the missing real path. No file was touched cosmetically while leaving a live bluff. (3 packages — `promptfoo/smolagents/superset` — used `inline fmt.Errorf` rather than a named sentinel; manually confirmed genuine.)

5. Mutation genuineness (step 4) — PASS (5/5, guards are non-tautologies)

Each: apply §1.1 mutation → standing guard FAILs → restore → guard PASSes.

#	Package	Mutation	Result
1	<code>gptengineer</code> (honest-error)	<code>generate</code> → <code>fabricated</code> <code>{files, status: generated}</code>	<code>TestD17_GPTEngineer_NoFabricatedFil</code> FAIL → restore ok

#	Package	Mutation	Result
2	claude_code (real-exec)	handleChat → "Claude Code received: %s" echo	TestPin_HandleChat_NotFabricatedEcho (sentinel marker absent) → restore ok
3	cursor (honest-error)	chat → fabricated {response}	TestCursor_Execute/chat... FAIL (expected got nil) → restore ok
4	cline (honest-error)	chat → fabricated {status:sent}	TestD17_Cline_NoFabrication FAIL → restore ok
5	smolagents (inline honest-error)	run → fabricated {result:Task completed successfully}	TestD17_Smolagents_RunIsHonestError → restore ok

Guards use unforgeable evidence: real-exec guards inject a fake binary via CLAUDE_BIN/GEMINI_BIN and assert its JSON-envelope sentinel flows through (impossible without genuine os/exec); honest-error guards assert require.Error + errors.Is(ErrCLINotWired) / “no headless CLI” substring. Not tautologies.

6. §11.4.120 reconciliations (step 5) — PASS

Spot-checked reconciled `_test.go` files assert the NEW mechanism, are not weakened to tautologies, and are not deleted-to-pass: - **cursor_test.go** — 6 cases flipped from `wantErr:false + assert.NotEmpty(m["response"/"code"/"status"])` to `wantErr:true, errMsg:"no headless CLI"`, with an explicit §11.4.120 comment. Harness enforces `require.Error + assert.Contains(err, errMsg)`. Session-CRUD cases stay genuine. - **gemini_test.go** — old "Generated by Gemini" assertions replaced by a sentinel-via-fake-binary real-exec proof + an absent-binary honest-error guard. - **promptfoo_test.go** — eval case reconciled to “honest error (no real run)” + dedicated `TestPromptfooEvalNoFabricatedScorecard` guard. - **claude_code_pin_test.go / cline_debluff_pin_test.go / gptengineer_stub_pin_test.go** — RED_MODE polarity-switch guards (§11.4.115): RED reproduces the pre-fix bluff, default GREEN asserts it is absent. - Full guard run (`- run 'TestD17|TestPin|NoFabricat|StubBluff|NotFabricated|NoFabricatedAIWork'`) across agents/...: **NO FAILURES**.

Findings

Must-fix (blocking this batch)

NONE.

Non-blocking / pre-existing (out of this batch’s scope — recommend follow-up tickets, do NOT block the D-16/D-17 commit)

1. `internal/clis/agents/claude_code/claude_code.go:477` — `handleMCP` action "list" returns a **hardcoded** server list "filesystem, github, memory, fetch, puppeteer" (BLUFF-002 class). Pre-existing; not in this diff.

2. `internal/clis/agents/claude_code/claude_code.go:489` — `handleMCP` action "call" returns `fmt.Sprintf("Called %s/%s via MCP", ...)` with no transport invoked (fabricated success). The sibling `mcp_integration.go` already defines `ErrMCPCliNotWired` for exactly this; `handleMCP` does not use it. Pre-existing; not in this diff.
3. `internal/clis/agents/agentdeck/agentdeck.go:197` and `internal/clis/agents/bridle/bridle.go:223,237` — "status": "completed" returned without a real run. Pre-existing; not in this diff.

These are genuine residual bluffs but lie outside the per-agent prompt→response handler scope this batch targeted; flagging them for a future de-bluff round.

VERDICT: GO

The D-16+D-17 de-bluff across the per-agent CLI layer is genuine, builds and tests cleanly as a whole layer, removes all live response-fabrication literals from the touched packages, and is backed by non-tautological standing regression guards (verified by 5/5 mutations) plus correct §11.4.120 test reconciliations. Scope is clean (agents/* only; no secrets; no mutation residue). The three non-blocking pre-existing findings are outside this batch.

Whole-internal/clis gate (steps 1–2)

- `go build ./internal/clis/...` → **PASS** (exit 0)
- `go vet ./internal/clis/...` → **PASS** (exit 0)
- `go test ./internal/clis/... -count=1` → **PASS** — 78 packages ok, **0 FAIL**, 4 [no test files] (pre-existing, untouched).

Finding counts

- Must-fix (blocking): **0**
- Non-blocking / pre-existing out-of-scope: **3**